ParaDG: An Exploratory Negative Study of Disagreement-Gated Paraphrase Blending for Text-to-Image Diffusion

Anonymous Authors

Abstract

Text-to-image diffusion models are sensitive to harmless prompt rewording: logically equivalent prompts can change attribute binding, relations, and object counts. We study whether this wording brittleness can be reduced at inference time by blending only checklist-audited paraphrases and increasing that blend only when the paraphrases disagree. Our method, ParaDG, is a training-free specialization of multi-prompt blending for Stable Diffusion v1.5. The core hypothesis is tested only against implemented paraphrase-based baselines: vanilla SD1.5, static paraphrase consensus, and adaptive ungated paraphrase blending. We evaluate these methods on 60 held-out prompts from attribute binding, relations, and numeracy, plus a 24-prompt robustness subset with paraphrase variants, using three seeds and official T2I-CompBench-family category metrics. The main result is negative. On held-out faithfulness, ParaDG is statistically indistinguishable from the implemented baselines: its highest mean category score is attribute binding 0.4677 (95% CI $[0.3629, 0.5723]$) versus vanilla SD1.5 at 0.4658 (95% CI $[0.3616, 0.5714]$). On the robustness subset, ParaDG underperforms adaptive ungated paraphrase blending on paraphrase robustness score, reaching 0.5509 (95% CI $[0.5015, 0.5973]$) versus 0.6111 (95% CI $[0.5617, 0.6636]$), while preserving similar LPIPS diversity (0.7650 versus 0.7620). The evidence therefore supports a narrower conclusion than the original proposal: paraphrase-conditioned consensus helps, but the present disagreement gate does not outperform simpler ungated blending under this evaluation setup. Because faithful AAPB reproduction, a strong single-prompt baseline, and human validation were not completed in this workspace, we present the paper as an exploratory negative-result study rather than a confirmatory method paper.

1 Introduction

Text-to-image diffusion models can produce materially different images for prompts that mean the same thing. Benchmarks such as T2I-CompBench and T2I-CompBench++ expose failures in attribute binding, spatial relations, and numeracy [Huang et al., 2023, 2025]. Structured Semantic Alignment and MetaLogic show that the problem is not limited to a single benchmark phrasing: semantically equivalent prompt rewrites themselves reveal robustness failures [Feng et al., 2023, Shen et al., 2026]. If harmless wording changes alter the generated scene, prompt-faithful generation remains brittle.

This paper tests one narrow hypothesis: *if disagreement among verified paraphrases is a useful signal of wording sensitivity, then a disagreement-gated paraphrase blend should outperform the implemented paraphrase-based baselines on robustness metrics, without causing diversity collapse*. We operationalize that hypothesis with ParaDG, a training-free extension of multi-prompt blending for Stable Diffusion v1.5. The method uses only audited paraphrases of the same prompt as auxiliary conditions and scales the consensus correction by the variance across paraphrase-conditioned predictions.

Preprint. Under review.

The completed experiments support a negative answer. On the 60-prompt faithfulness split, ParaDG is competitive but its small gains fall within overlapping 95% confidence intervals. On the 24-prompt robustness subset, adaptive ungated paraphrase blending is stronger on paraphrase robustness score (PRS), DINO consistency, and CLIP image consistency. The same pattern appears in ablations: removing the gate does not hurt and slightly improves attribute binding on the matched analysis subset. Together, these results suggest that paraphrase-conditioned consensus is the helpful ingredient, while the current variance gate is not.

Our contributions are:

- We instantiate and test ParaDG, a training-free paraphrase-restricted diffusion update that combines sparse multi-prompt consensus with a variance-based disagreement gate and slot-aware token redistribution.
- We evaluate the method on a controlled prompt matrix with 72 prompts, 144 accepted paraphrases, official T2I-CompBench-family category metrics, and 95% confidence intervals reported directly in the main tables.
- We report an explicit negative result: under the available automatic evaluation pipeline and against the implemented paraphrase-based baselines, disagreement gating does not improve over simpler ungated paraphrase blending.
- We audit and surface the incomplete parts of the original plan in the main paper so that the contribution is the negative finding itself, not an implicit claim of experimental completeness.

Section 2 reviews related work, Section 3 defines ParaDG and the tested hypothesis, Section 4 presents the experiments, Section 5 discusses what likely failed, and Appendix A summarizes reproducibility artifacts and missing components.

2 Related Work

Diffusion models and compositional evaluation. Our experiments use Stable Diffusion v1.5, built on latent diffusion modeling [Rombach et al., 2022]. T2I-CompBench and T2I-CompBench++ provide the main compositional framing for attribute binding, relations, and numeracy [Huang et al., 2023, 2025].

Prompt robustness evaluation. Structured Semantic Alignment stresses text-to-image systems with controlled semantic perturbations [Feng et al., 2023]. MetaLogic evaluates whether logically equivalent prompts produce semantically consistent images [Shen et al., 2026]. Our target is closest to these works, but we consume equivalent prompts at inference time rather than only as an evaluator.

Training-free guidance for prompt faithfulness. Attend-and-Excite, A-STAR, MaskDiffusion, and Self-Coherence Guidance all improve alignment from a single prompt [Chefer et al., 2023, Agarwal et al., 2023, Zhou et al., 2023, Wang et al., 2025]. These methods already establish test-time semantic guidance, so our contribution is not a new guidance family.

Multi-prompt blending and prompt optimization. The closest method is Adaptive Auxiliary Prompt Blending (AAPB), which derives an adaptive coefficient for blending a target prompt with an auxiliary prompt [Lee et al., 2026]. ParaDG differs in two ways: it restricts auxiliary prompts to audited paraphrases of the same scene, and it activates blending according to observed disagreement across those paraphrases. Prompt optimization systems such as NeuroPrompts seek stronger prompt wordings [Rosenman et al., 2024]; our goal is instead to make semantically equivalent wordings behave consistently.

Unlike prior evaluation-only robustness work, ParaDG turns equivalent-prompt disagreement into an inference-time signal. Unlike generic prompt blending, it asks whether that signal is scientifically useful beyond ungated paraphrase consensus.

3 Method

3.1 Prompt Schema and Paraphrase Audit

Each original prompt p_0 is represented with a fixed slot schema

$$\{\text{count}_1, \text{object}_1, \text{attribute}_1, \text{relation}, \\ \text{count}_2, \text{object}_2, \text{attribute}_2\}. \tag{1}$$

where fields may be empty but every prompt must be representable with at most two objects. This deliberately narrows the study to categories where Stable Diffusion v1.5 commonly fails and where semantic equivalence can be audited with low ambiguity.

Paraphrases are not generated freely. They are produced by controlled rewrite templates: attribute-order swaps, existential rewrites, clause-order rewrites, fronted-relation rewrites, and count-preserving digit rewrites. A candidate is accepted only if every populated slot is preserved exactly, no new object or relation is introduced, and directed relations remain unchanged.

The completed audit artifact contains 72 prompts in total: 12 pilot prompts and 60 held-out prompts. It generated 148 candidate paraphrases, accepted 144, and rejected 4 for number-grammar mismatch. All 60 held-out prompts retained two approved paraphrases, giving full overlap coverage across attribute binding, relations, and numeracy. A 48-pair secondary audit reports Cohen's $\kappa = 1.0$, but this was an automatic-primary versus automatic-secondary check rather than the preregistered human double annotation, so we use it only as an internal consistency check.

3.2 Disagreement-Gated Paraphrase Blending

Let $P = \{p_0, p_1, \ldots, p_m\}$ be the approved paraphrase set for an original prompt p_0. At a guided diffusion step t, the conditional denoiser prediction for prompt p_j is $\epsilon_t(p_j)$. We form the paraphrase consensus

$$\bar{\epsilon}_t = \frac{1}{|P|} \sum_{j=0}^{m} \epsilon_t(p_j)$$

and update the original-prompt prediction as

$$\epsilon_t^{\text{ParaDG}} = \epsilon_t(p_0) + \alpha_t g_t \left(\bar{\epsilon}_t - \epsilon_t(p_0)\right).$$

Here α_t is a fixed step-dependent blend schedule and g_t is a disagreement gate,

$$g_t = \min\left(1,\ c \left\|\text{Var}_j[\epsilon_t(p_j)]\right\|_1\right).$$

The final configuration, selected on the pilot by runtime among tied candidates, uses $c = 0.5$, guided steps $[2, 4, 6, 9, 12, 16, 20, 24]$, and $\alpha_t = [0.6, 0.6, 0.5, 0.5, 0.4, 0.3, 0.2, 0.2]$.

3.3 Slot-Aware Redistribution

The correction term $\bar{\epsilon}_t - \epsilon_t(p_0)$ is redistributed over token groups aligned to the slot schema. For each populated slot, we aggregate cross-attention over the tokens that realize that slot in each paraphrase, estimate slot disagreement, and use the normalized disagreement weights to apportion the already-computed correction. This adds no learned parameters and no external parser.

3.4 Tested Hypothesis

The method is intentionally evaluated against a stronger null than vanilla generation. The scientific question is not whether paraphrase information helps at all, but whether *disagreement gating* adds value beyond ungated paraphrase blending. The key hypothesis is:

> If paraphrase disagreement tracks wording brittleness, then ParaDG should outperform both static consensus and adaptive ungated blending on robustness metrics such as PRS and DINO consistency, without a severe LPIPS drop relative to vanilla SD1.5.

The experiments below were designed to falsify that claim if paraphrase consensus alone explained the gains.

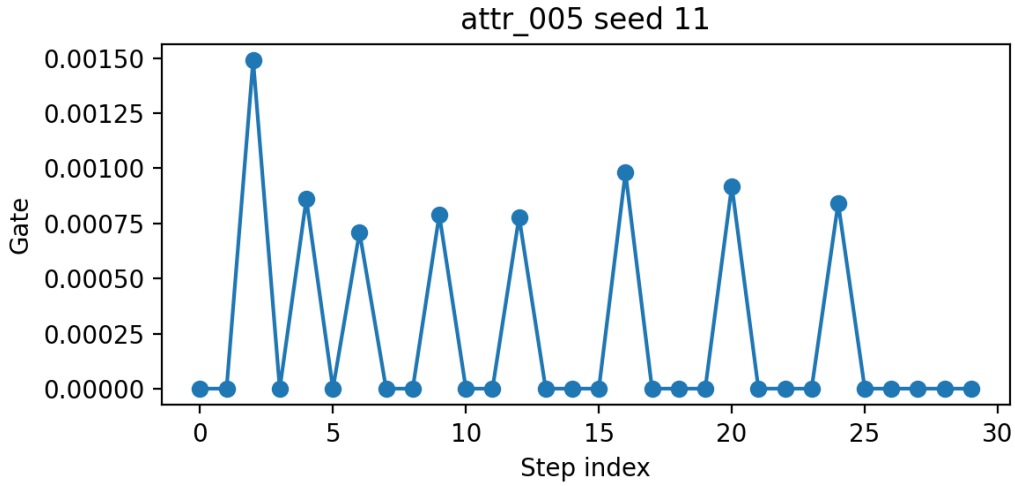

Figure 1: Example gate trace for one held-out attribute-binding prompt. The gate is active only on the eight designated guided steps and zero elsewhere. Source data are stored in figures/gate_trace_1.json.

4 Experiments

4.1 Setup

Backbone and generation. All methods use the runwayml/stable-diffusion-v1-5 checkpoint with DDIM sampling, 30 denoising steps, CFG scale 7.5, resolution 512×512, and seeds $\{11, 22, 33\}$. The experiments ran in the recorded environment from exp/environment_setup/run_manifest.json on a single RTX A6000 48GB machine.

Prompt matrix. The main faithfulness study uses 60 held-out prompts, evenly split across attribute binding, relations, and numeracy. The robustness study uses 24 overlap prompts with paraphrase variants. With three seeds, this yields 180 generations per method for faithfulness and 216 generations per method for robustness.

Methods. We compare four implemented methods: vanilla SD1.5, static paraphrase consensus, adaptive ungated paraphrase blending, and ParaDG. These are the only methods used to test the core hypothesis. The preregistered faithful AAPB reproduction, the preregistered single-prompt baseline, the human reliability subset, and the human double annotation were all formally revised to infeasible in this workspace, so we report them as missing constraints rather than silently replacing them.

Metrics. Primary semantic metrics are the official T2I-CompBench-family category scores for attribute binding, relations, and numeracy. Robustness diagnostics additionally include PRS, CLIP image consistency, DINO consistency, robustness-split CLIPScore, LPIPS seed diversity, DINO seed dispersion, ImageReward, and latency. Because the human reliability subset was not completed, the robustness metrics remain automatic diagnostics rather than human-validated conclusions.

4.2 Paraphrase Audit and Pilot

Table 2 summarizes the data preparation artifact. Table 1 makes the paper's scope explicit by naming the revised and skipped result files directly: the completed study is substantial enough to support an audited exploratory result, but not enough to support the stronger confirmatory claims originally planned. The pilot selected the final setting among nine gate-schedule candidates; all nine tied on mean pilot category score at 0.2083, so the final setting was chosen by runtime. The projected full-study cost was 1.3836 GPU hours, and the reduced matrix was not activated.

Table 1: Artifact-status summary for the plan components that most affect interpretability. The table names the exact result files so the completed, revised, and skipped pieces are directly traceable from the paper text.

Artifact file	Role	Status
exp/data_preparation/results.json	Prompt split and paraphrase audit	Completed
exp/pilot/results.json	Pilot schedule selection	Completed
exp/official_compbench_eval/results.json	Official metric backfill log	Completed
exp/metric_backfill/results.json	Row-level metric propagation log	Completed
exp/aapb_reproduction/results.json	AAPB reproduction attempt	Revised infeasible
exp/single_prompt_baseline/results.json	Single-prompt baseline attempt	Revised infeasible
exp/human_reliability_subset/results.json	Human robustness subset	Revised infeasible
exp/paraphrase_double_annotation/results.json	Human double annotation	Revised infeasible
exp/plan_revision/results.json	Formal revision record	Completed

Table 2: Paraphrase-audit and pilot summary from exp/data_preparation/results.json and exp/pilot/results.json.

Statistic	Value
Total prompts	72
Held-out prompts	60
Overlap prompts with two paraphrases	60
Candidate paraphrases	148
Accepted paraphrases	144
Rejected paraphrases	4
Secondary audit pairs	48
Secondary audit κ	1.0
Best gate constant c	0.5
Best schedule	early-heavy
Projected GPU hours	1.3836

4.3 Main Faithfulness Results

Table 3 reports the main held-out results with 95% confidence intervals. The scientific signal is weak. ParaDG has the highest mean on attribute binding, 0.4677, but its interval almost exactly overlaps vanilla SD1.5 (0.4658) and static consensus (0.4615). Static consensus is best on relations and numeracy. Put differently, the largest faithfulness differences are small relative to uncertainty, so the faithfulness results do not support a claim that disagreement gating meaningfully helps.

4.4 Robustness Results

Table 4 makes the main negative result explicit. Adaptive ungated blending is best on PRS, DINO consistency, and CLIP image consistency, while ParaDG is only marginally higher than vanilla SD1.5 on PRS and essentially tied with vanilla on DINO consistency. ParaDG does achieve the highest robustness-split CLIPScore, but the difference is small and does not translate into better robustness under equivalent prompt rewrites.

The robustness table is also inconsistent with a diversity-collapse explanation. LPIPS remains in a narrow range from 0.7620 to 0.7652, and ParaDG nearly matches vanilla SD1.5 on that metric. The more plausible interpretation is that paraphrase consensus helps, but the gate itself neither calibrates that signal well nor applies a sufficiently strong correction.

Table 3: Held-out faithfulness results on 60 prompts and 180 generations per method. Values are mean [95% CI]. Best means are in **bold**. The near-complete CI overlap indicates that the small deltas should be treated as within noise.

Method	Attr. bind. ↑	Relations ↑	Numeracy ↑	CLIPScore ↑	Latency (s) ↓
Vanilla SD1.5	0.4658 [0.3616, 0.5714]	0.1622 [0.0802, 0.2567]	0.5750 [0.5000, 0.6417]	0.2632 [0.2582, 0.2685]	**1.9035** [1.8949, 1.9132]
Static consensus	0.4615 [0.3631, 0.5681]	**0.1850** [0.1024, 0.2818]	**0.6083** [0.5333, 0.6833]	0.2649 [0.2595, 0.2701]	2.7885 [2.7811, 2.7965]
Adaptive ungated	0.4515 [0.3491, 0.5566]	0.1647 [0.0831, 0.2562]	0.5833 [0.5167, 0.6500]	**0.2652** [0.2596, 0.2705]	2.7938 [2.7873, 2.8009]
ParaDG	**0.4677** [0.3629, 0.5723]	0.1757 [0.0928, 0.2730]	0.5667 [0.4917, 0.6417]	0.2631 [0.2581, 0.2683]	3.2771 [3.2677, 3.2871]

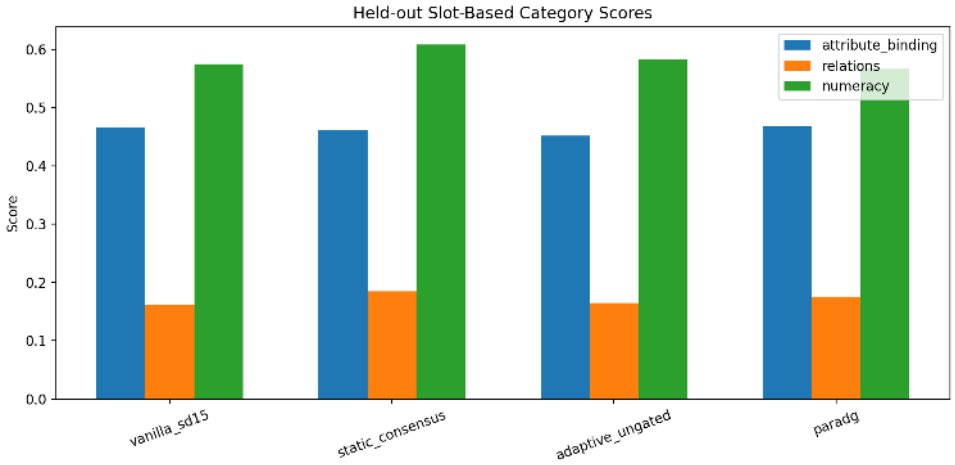

Figure 2: Category-wise faithfulness summary. ParaDG has the highest mean on attribute binding, but static consensus is stronger on relations and numeracy. The source tables are in figures/category_scores_source.json and figures/category_scores_source.csv.

4.5 Ablations and Failure Analysis

Table 5 compares ParaDG against matched ablations. The pattern is consistent with the main robustness table: removing the gate slightly raises the matched-slice attribute-binding mean from 0.5567 to 0.5724, and all confidence intervals overlap heavily. Removing slot redistribution or timestep adaptation barely changes the semantic scores. On the separate non-equivalent comparison, the matched ParaDG slice reaches attribute binding 0.4870 while the non-equivalent auxiliary-prompt ablation reaches 0.4774, suggesting that semantic equivalence still matters even though the present gate does not.

The gate traces in Figure 1 help explain why. The gate activates at the intended steps, but the released traces show magnitudes on the order of 10^{-3}. Combined with the flat pilot, this suggests two likely failure modes: the disagreement statistic is too weak to separate useful from noisy paraphrase variance, and the resulting correction is too small relative to the already helpful ungated consensus term.

5 Discussion and Limitations

The main scientific conclusion is narrower than the proposal: *equivalent-prompt consensus helps, but the current disagreement gate does not*. The evidence for that statement is consistent across the paper. On faithfulness, the observed gains are tiny and within overlapping confidence intervals. On robustness, ungated blending is better than ParaDG on the main automatic robustness metrics. On ablations, removing the gate does not hurt. And on diagnostics, the pilot is flat and the gate magnitudes are small.

The likely reason is not that paraphrases are useless. In fact, the stronger implemented baseline is precisely the ungated paraphrase blend. A more plausible explanation is that the variance-only gate is a poor proxy for harmful wording sensitivity. It can be triggered by benign prompt-conditioned

Table 4: Robustness results on the 24-prompt overlap subset and 216 generations per method. Values are mean [95% CI]. Best means are in **bold**.

Method	Attr. bind. ↑	Relations ↑	Numeracy ↑	PRS ↑	DINO cons. ↑	CLIP img. cons. ↑	CLIPScore ↑	LPIPS ↑	Latency (s) ↓
Vanilla SD1.5	0.5616 [0.4125, 0.7166]	0.0682 [0.0000, 0.1780]	0.6250 [0.5203, 0.7292]	0.5370 [0.4876, 0.5849]	0.6538 [0.6291, 0.6800]	0.8202 [0.8063, 0.8333]	0.2586 [0.2532, 0.2640]	**0.7652** [0.7452, 0.7866]	**1.8887** [1.8817, 1.8959]
Static consensus	**0.5981** [0.4498, 0.7407]	**0.1336** [0.0249, 0.2671]	0.6667 [0.5417, 0.7917]	0.5741 [0.5216, 0.6296]	0.6733 [0.6495, 0.6982]	0.8300 [0.8165, 0.8427]	0.2576 [0.2523, 0.2630]	0.7640 [0.7443, 0.7852]	2.7891 [2.7829, 2.7954]
Adaptive ungated	0.5862 [0.4408, 0.7265]	0.0899 [0.0000, 0.2031]	**0.6875** [0.5833, 0.7917]	**0.6111** [0.5617, 0.6636]	**0.6875** [0.6632, 0.7124]	**0.8389** [0.8251, 0.8515]	0.2577 [0.2523, 0.2630]	0.7620 [0.7423, 0.7834]	2.7998 [2.7920, 2.8098]
ParaDG	0.5656 [0.4150, 0.7195]	0.0685 [0.0000, 0.1785]	0.6458 [0.5208, 0.7500]	0.5509 [0.5015, 0.5973]	0.6532 [0.6281, 0.6792]	0.8206 [0.8067, 0.8338]	**0.2587** [0.2532, 0.2641]	0.7650 [0.7450, 0.7862]	3.2797 [3.2710, 3.2881]

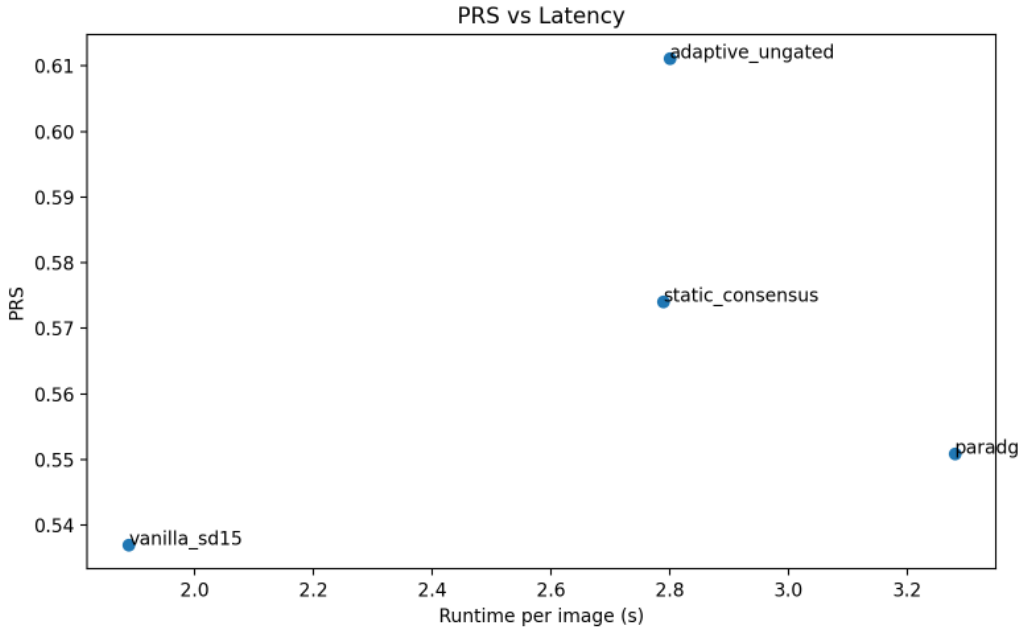

Figure 3: PRS versus latency on the robustness subset. Adaptive ungated blending offers the best automatic robustness among implemented methods, while vanilla SD1.5 remains cheapest. Source data are in figures/prs_vs_latency_source.json and figures/prs_vs_latency_source.csv.

variance, and when it is triggered its correction remains weak relative to the baseline consensus update.

Compact error analysis. The qualitative grid and matched ablations suggest three recurring failure patterns. First, relation prompts remain brittle across all methods, with robustness-split relations scores of only 0.0682 for vanilla SD1.5, 0.0899 for adaptive ungated blending, and 0.0685 for ParaDG, indicating that paraphrase consensus alone does not reliably repair directed spatial structure. Second, numeracy often improves under consensus-style methods, but ParaDG does not preserve that gain: on the robustness subset adaptive ungated blending reaches 0.6875 numeracy while ParaDG reaches 0.6458. Third, the gate traces show activation at the intended steps but with small magnitudes, which is consistent with a mechanism that perturbs the denoiser too weakly to alter difficult failure cases.

Several limitations materially constrain the paper. First, the faithful AAPB reproduction and the strong single-prompt baseline were not completed, so this study cannot support a strong comparative claim against the closest alternative families. Second, the human reliability subset and human double annotation were skipped, so the robustness story remains exploratory rather than confirmatory. Third, numeracy uses the documented RS200 compatibility fallback because the originally referenced R50 checkpoint was unavailable. Finally, the slot schema is intentionally narrow and does not test broader natural-language prompt variation.

Those limitations matter enough that, in its current state, this paper is better positioned as an exploratory negative-results submission than as a full confirmatory method paper. The next round of work should prioritize the missing strong baselines and human validation before introducing a more sophisticated gating mechanism.

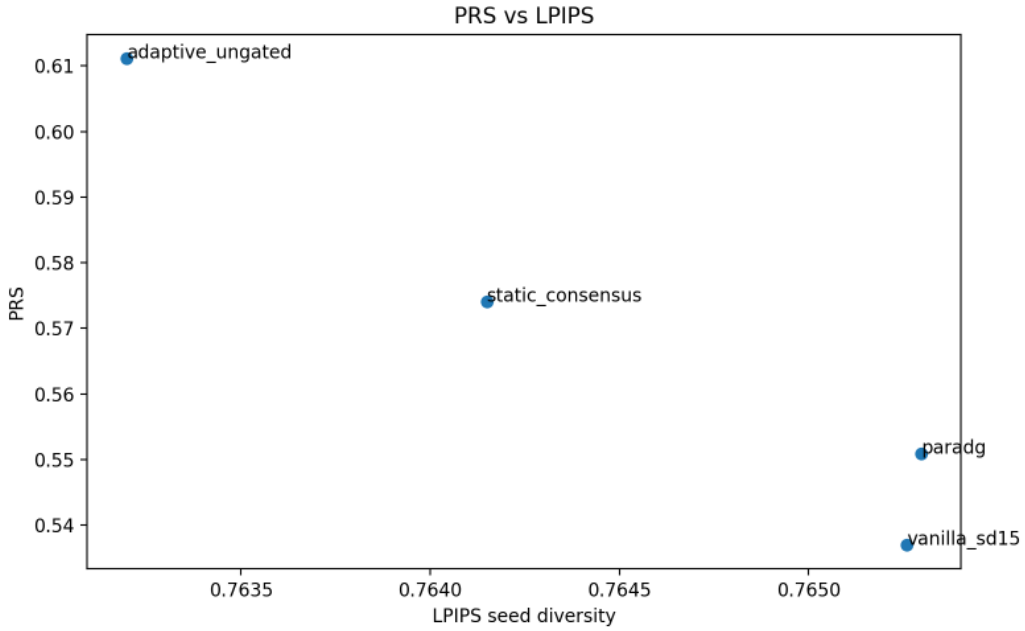

Figure 4: PRS versus LPIPS seed diversity on the robustness subset. The methods occupy a narrow LPIPS band, so the lower PRS of ParaDG is not obviously caused by diversity collapse. Source data are in figures/prs_vs_lpips_source.json and figures/prs_vs_lpips_source.csv.

Table 5: Ablations on the 18-prompt matched analysis subset. Values are mean [95% CI]. Best means are in **bold**.

Method	Attr. bind. ↑	Relations ↑	Numeracy ↑	CLIPScore ↑	Latency (s) ↓
ParaDG (matched subset)	0.5567 [0.4342, 0.6792]	0.0913 [0.0036, 0.1790]	0.6111 [0.5070, 0.7152]	0.2691 [0.2617, 0.2765]	3.2776 [3.2640, 3.2912]
No gate	**0.5724** [0.4021, 0.7340]	0.0908 [0.0000, 0.2370]	0.6111 [0.4722, 0.7500]	0.2694 [0.2587, 0.2796]	3.2819 [3.2593, 3.3064]
No slot redistribution	0.5572 [0.3853, 0.7284]	0.0906 [0.0000, 0.2367]	0.6111 [0.4722, 0.7500]	**0.2696** [0.2594, 0.2795]	**2.7870** [2.7738, 2.8034]
No timestep adaptation	0.5506 [0.3770, 0.7225]	0.0912 [0.0000, 0.2379]	0.6111 [0.4722, 0.7500]	0.2693 [0.2588, 0.2792]	3.2699 [3.2501, 3.2938]
One paraphrase only	0.5630 [0.3987, 0.7282]	**0.0915** [0.0000, 0.2384]	0.6111 [0.4722, 0.7500]	0.2692 [0.2589, 0.2791]	2.8148 [2.7978, 2.8338]

6 Conclusion

We presented ParaDG, a training-free paraphrase-restricted diffusion update that uses audited paraphrases and a variance-based disagreement gate to stabilize generation under equivalent prompt rewrites. Across 60 held-out prompts and a 24-prompt robustness subset, the results do not support the intended hypothesis. The best implemented robustness baseline is adaptive ungated blending, which reaches PRS 0.6111 versus 0.5509 for ParaDG, while faithfulness differences remain within overlapping confidence intervals.

This is still an informative result. It suggests that paraphrase-conditioned consensus is a real inference-time signal for text-to-image generation, but that the present disagreement gate is not the right way to exploit it. A stronger follow-up should first complete faithful AAPB and single-prompt baselines, plus human robustness validation, and only then revisit richer gating rules or learned paraphrase weighting.

A Reproducibility Summary

Every reported number in the paper comes from the stored experiment summaries under exp/. The main result files are exp/vanilla_sd15/results.json, exp/static_consensus/results.json, exp/adaptive_ungated/results.json, and exp/paradg/results.json. The matched ablation summaries are exp/paradg_matched_subset/results.json, exp/ablation_no_gate/results.json, exp/ablation_no_slot/results.json, exp/ablation_no_timestep/results.json, exp/ablation_reduced_paraphrase/results.json, exp/ablation_nonequivalent/results.json,

Table 6: Non-equivalent auxiliary-prompt comparison on the 12-prompt analysis slice. The first row is the matched ParaDG slice from exp/paradg_nonequivalent_matched/results.json; the second is the dedicated non-equivalent rerun from exp/ablation_nonequivalent/results.json. Values are mean [95% CI].

Method	Attr. bind. ↑	Relations ↑	Numeracy ↑
ParaDG (equivalent prompts)	**0.4870** [0.3286, 0.6453]	**0.0538** [-0.0191, 0.1268]	0.5417 [0.4109, 0.6725]
Non-equivalent auxiliary prompt	0.4774 [0.2598, 0.6810]	0.0531 [0.0000, 0.1592]	0.5417 [0.3750, 0.7083]

and exp/paradg_nonequivalent_matched/results.json. Figure source data are released under figures/.

References

Aishwarya Agarwal, Srikrishna Karanam, K. J. Joseph, Apoorv Saxena, Koustava Goswami, and Balaji Vasan Srinivasan. A-star: Test-time attention segregation and retention for text-to-image synthesis. *arXiv preprint arXiv:2306.14544*, 2023. URL https://arxiv.org/abs/2306.14544.

Hila Chefer, Yuval Alaluf, Yael Vinker, Lior Wolf, and Daniel Cohen-Or. Attend-and-excite: Attention-based semantic guidance for text-to-image diffusion models. *ACM Transactions on Graphics*, 2023. URL https://arxiv.org/abs/2301.13826.

Qianyu Feng, Yulei Sui, and Hongyu Zhang. Uncovering limitations in text-to-image generation: A contrastive approach with structured semantic alignment. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 8876–8888, 2023. URL https://aclanthology.org/2023.findings-emnlp.595/.

Kaiyi Huang, Kaiyue Sun, Enze Xie, Zhenguo Li, and Xihui Liu. T2i-compbench: A comprehensive benchmark for open-world compositional text-to-image generation. In *Advances in Neural Information Processing Systems*, volume 36, pages 78723–78747, 2023. URL https://papers.neurips.cc/paper_files/paper/2023/hash/f8ad010cdd9143dbb0e9308c093aff24-Paper-Datasets_and_Benchmarks.pdf.

Kaiyi Huang, Chengqi Duan, Kaiyue Sun, Enze Xie, Zhenguo Li, and Xihui Liu. T2i-compbench++: An enhanced and comprehensive benchmark for compositional text-to-image generation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2025. URL https://arxiv.org/abs/2307.06350.

Kwanyoung Lee, SeungJu Cha, Yebin Ahn, Hyunwoo Oh, Sungho Koh, and Dong-Jin Kim. Adaptive auxiliary prompt blending for target-faithful diffusion generation. *arXiv preprint arXiv:2603.19158*, 2026. URL https://arxiv.org/abs/2603.19158. Accepted to CVPR 2026 main track.

Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Bjorn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022. URL https://arxiv.org/abs/2112.10752.

Shachar Rosenman, Vasudev Lal, and Phillip Howard. Neuroprompts: An adaptive framework to optimize prompts for text-to-image generation. In *Proceedings of the EACL 2024 System Demonstration Track*, 2024. URL https://arxiv.org/abs/2311.12229.

Yifan Shen, Yangyang Shu, Hye-young Paik, and Yulei Sui. Metalogic: Robustness evaluation of text-to-image models via logically equivalent prompts. In *Formal Methods and Software Engineering: 26th International Conference on Formal Engineering Methods (ICFEM 2025)*, pages 152–171, 2026. URL https://arxiv.org/abs/2510.00796. Conference held in November 2025; arXiv:2510.00796.

Shulei Wang, Wang Lin, Hai Huang, Hanting Wang, Sihang Cai, WenKang Han, Tao Jin, Jingyuan Chen, Jiacheng Sun, Jieming Zhu, and Zhou Zhao. Towards transformer-based aligned generation with self-coherence guidance. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025. URL https://arxiv.org/abs/2503.17675.

Yupeng Zhou, Daquan Zhou, Zuo-Liang Zhu, Yaxing Wang, Qibin Hou, and Jiashi Feng. Maskdiffusion: Boosting text-to-image consistency with conditional mask. *arXiv preprint arXiv:2309.04399*, 2023. URL https://arxiv.org/abs/2309.04399.

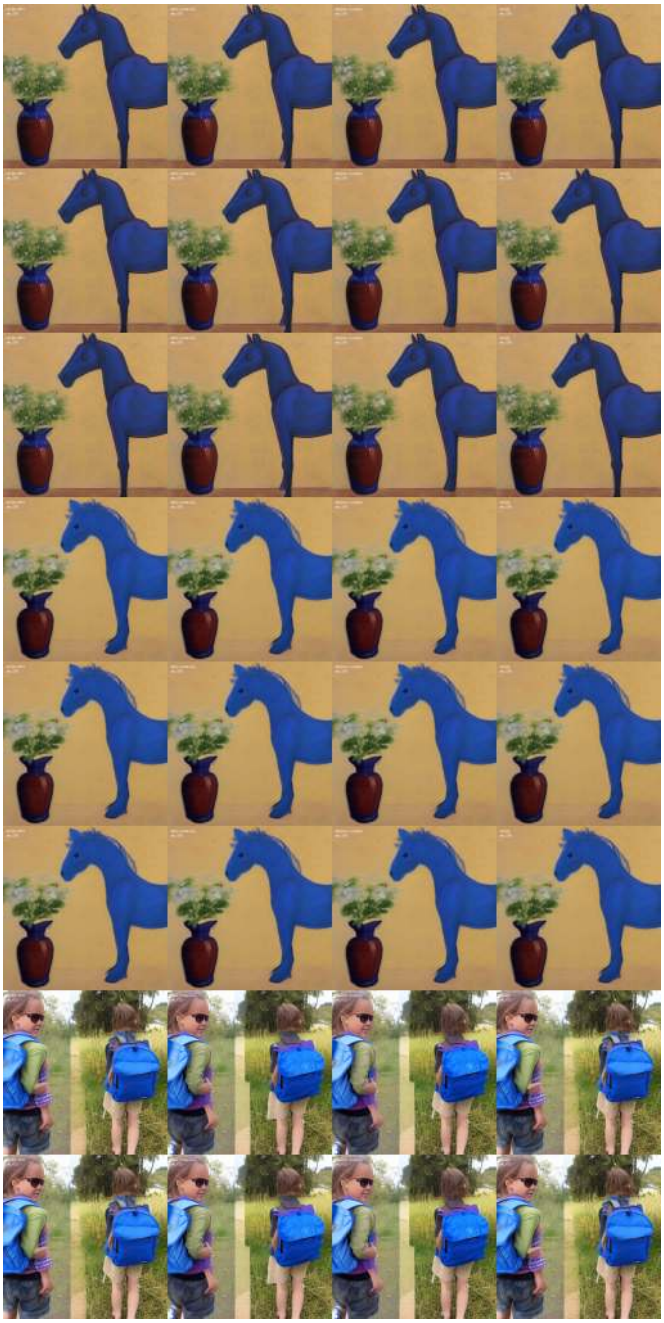

Figure 5: Qualitative grid assembled automatically from the first eight robustness prompts encountered in exp/paradg/generation_index.csv: `attr_005` through `attr_012`. Each row is one original prompt at seed 11, and the four columns are Vanilla SD1.5, static consensus, adaptive ungated blending, and ParaDG, in that order. Prompt IDs and method names are overlaid directly on the cells by the figure-generation script, so this is an uncurated attribute-binding slice rather than a balanced cross-category showcase. The figure manifest is stored in figures/qualitative_grid_manifest.json.