On the Challenges of Multi-Fidelity Conditional Independence Testing for Causal Discovery: An Empirical Study

Anonymous Author(s)
Anonymous Institution
anonymous@email.com

April 8, 2026

Abstract

Multi-fidelity methods have achieved significant success in Bayesian optimization by strategically balancing cheap approximations against expensive accurate evaluations. Motivated by this success, we investigate whether similar principles can be applied to constraint-based causal discovery, where conditional independence (CI) testing dominates computational cost. We propose Multi-Fidelity Adaptive Causal Discovery (MF-ACD), a framework that uses tiered CI tests—correlation-based screening (low-fidelity), Fisher's Z-test (medium-fidelity), and kernel-based tests (high-fidelity)—with an adaptive Uncertainty-Guided Fidelity Selection (UGFS) mechanism. Through extensive experiments on synthetic graphs with 20–50 nodes, we find that while MF-ACD reduces proxy cost metrics by 49%, it suffers from 34% lower F1 score compared to a simple correlation-only baseline (0.281 vs. 0.425) and requires 26× more wall-clock time (4.98s vs. 0.19s). Furthermore, our adaptive allocation mechanism shows no statistically significant improvement over fixed allocation ($p = 0.847$). These negative results reveal fundamental challenges: unlike Bayesian optimization where cheap approximations capture continuous function structure, CI tests make hard binary decisions with irreversible information loss, and the overhead of adaptive mechanisms negates theoretical cost savings.

1 Introduction

Causal discovery—the task of inferring causal relationships from observational data—is a fundamental problem in science and engineering. Constraint-based methods, particularly the PC algorithm spirtes2000causation, are among the most widely used approaches due to their theoretical guarantees and interpretability. However, these methods face a critical scalability bottleneck: the number of conditional independence (CI) tests required grows factorially with the number of variables, and each test can be computationally expensive, especially for non-linear dependencies or high-dimensional conditioning sets.

Current approaches to address this scalability challenge fall into three categories: (1) hardware acceleration zarebavani2019cupc requiring specialized hardware, (2) algorithmic heuristics that approximate the original algorithm often sacrificing accuracy, and (3) divide-and-conquer methods dong2025dcilp, shanmugam2021hccd that partition variables but use fixed heuristics for merging. These approaches share a fundamental limitation: they treat all CI tests equally, ignoring that different tests have different computational costs and information values.

Multi-fidelity methods have achieved remarkable success in Bayesian optimization kandasamy2016gaussian, kandasamy2017multi by strategically balancing cheap approximations against expensive accurate evaluations. This raises a natural question: **Can multi-fidelity principles be applied to CI testing in causal discovery to reduce computational cost while maintaining accuracy?**

1.1 Our Approach

We propose Multi-Fidelity Adaptive Causal Discovery (MF-ACD), a three-phase framework:

1. **Phase 1 (Low-Fidelity):** Correlation-based screening with FDR control
2. **Phase 2 (Medium-Fidelity):** Fisher's Z-test on candidate edges
3. **Phase 3 (High-Fidelity):** Kernel-based tests for critical edges

The key innovation is our Uncertainty-Guided Fidelity Selection (UGFS) mechanism that adaptively allocates the computational budget based on expected information gain per unit cost.

1.2 Contributions

Our contributions are:

- We propose MF-ACD, the first systematic investigation of multi-fidelity CI testing for causal discovery with an adaptive information-theoretic allocation mechanism.
- We conduct extensive experiments on synthetic benchmarks, reporting results with 95% confidence intervals from 360 total runs using fixed random seeds.
- We present a detailed analysis of **why multi-fidelity approaches fail for CI testing**: irreversible information loss from binary test decisions, high correlation between fidelity levels, and adaptive overhead that negates theoretical savings.
- We provide insights that may guide future research toward more promising directions for scalable causal discovery.

2 Related Work

Constraint-Based Methods. The PC algorithm spirtes2000causation is the foundation of constraint-based causal discovery. PC-Stable colombo2014order improves order-independence but does not address computational cost. C-PC lee2025constraint uses restricted conditioning sets to improve reliability. These methods use a uniform CI test throughout, unlike our multi-fidelity approach.

Hierarchical and Divide-and-Conquer Methods. HCCD shanmugam2021hccd recursively clusters variables using normalized min-cut and applies PC in bottom-up fashion. DCILP dong2025dcilp uses Markov blankets for partitioning and ILP for merging local subgraphs. VISTA shah2024vista uses voting-based integration of subgraphs. These methods reduce the number of CI tests through clustering, whereas MF-ACD attempts to reduce cost through adaptive fidelity selection.

Significance-Weighted Approaches. SWCD bai2025significance uses Path Significance Values (PSV) for graph partitioning. While SWCD applies significance to structure decisions (partitioning), MF-ACD applies information-theoretic selection to statistical test decisions (fidelity).

Multi-Fidelity Methods. Multi-fidelity Bayesian optimization kandasamy2016gaussian, kandasamy2017multi uses cheap approximations to guide expensive evaluations. In causal discovery, zhang2023bayesian address active causal discovery with interventions (not observational data). MF-ACD differs in targeting observational causal discovery with tiered test fidelities.

Tiered Background Knowledge. bang2023wise, bang2025constraint introduced tiered background knowledge—exploiting temporal ordering to improve identifiability. While they use tiered variable structure as external knowledge, MF-ACD uses tiered test fidelity as an internal algorithmic choice.

3 Method

3.1 Overview: MF-ACD Framework

MF-ACD operates in three phases with adaptive budget allocation:

Phase 1: Low-Fidelity Skeleton Screening. We use Pearson and partial correlation tests with relaxed thresholds and Benjamini-Hochberg FDR control at $\alpha_1 = 0.10$. This phase processes $O(p^2)$ edges at low cost to identify candidate edges.

Phase 2: Medium-Fidelity Local Refinement. We apply Fisher's Z-tests on remaining candidate edges with Storey's adaptive FDR at $\alpha_2 = 0.05$, focusing on local neighborhoods.

Phase 3: High-Fidelity Critical Resolution. We apply kernel-based tests (KCI) only for edges with conflicting local orientations, using Holm-Bonferroni correction at $\alpha_3 = 0.01$.

3.2 Uncertainty-Guided Fidelity Selection (UGFS)

The UGFS mechanism adaptively selects edge-fidelity pairs based on expected information gain per unit cost:

Definition 1 (Information Gain)**.** *Let G denote the graph structure random variable and $X_{e,f} \in \{0, 1\}$ the outcome of CI test for edge e at fidelity f. The expected information gain is:*

$$IG(e, f) = H(P(G \mid D_{obs})) - \mathbb{E}_{X_{e,f}}[H(P(G \mid D_{obs}, X_{e,f}))] \tag{1}$$

where $H(\cdot)$ denotes the Shannon entropy.

Definition 2 (Fidelity Selection Score)**.** *The selection score for edge-fidelity pair (e, f) is:*

$$Score(e, f) = \frac{IG(e, f) \times SI(e)}{Cost(f)} \tag{2}$$

where $SI(e)$ is the structural importance (normalized Markov blanket size) and $Cost(f)$ is the normalized computational cost.

We approximate information gain as:

$$\mathrm{IG}(e, f) \approx 4p_e(1 - p_e) \cdot \mathrm{Power}(f) \cdot \mathrm{Disc}(f) \tag{3}$$

where p_e is the current belief that edge e exists, and $\mathrm{Power}(f)$, $\mathrm{Disc}(f)$ are the test power and discriminative power.

3.3 Budget Allocation

Based on pilot experiments measuring actual test costs across 24 configurations, we use initial allocations of approximately 30% low-fidelity, 20% medium-fidelity, and 50% high-fidelity. During execution, we adapt based on observed statistics using a regret minimization criterion with adaptation rate $\eta = 0.1$.

3.4 Theoretical Foundations

Theorem 1 (Soundness)**.** *Under the faithfulness assumption, if all CI tests at all fidelities are sound and multiple testing corrections are applied as specified, MF-ACD controls the family-wise error rate at level $\alpha_{total} \leq \alpha_1 + \alpha_2 + \alpha_3$.*

Proof Sketch. By the union bound, the probability of any false positive across all phases is bounded by the sum of phase-wise error rates. The sequential nature ensures that errors from early phases can be corrected in later phases. □

Theorem 2 (Budget Scaling)**.** *For sparse graphs with maximum degree d, let k be the number of edges receiving high-fidelity testing. The expected computational cost is:*

$$\mathbb{E}[C_{MF\text{-}ACD}] = O(pd^2 \cdot c_M + k \cdot c_H) \tag{4}$$

compared to $O(p^2 \cdot c_H)$ for full high-fidelity discovery.

Table 1: Main results comparing MF-ACD against baselines on synthetic graphs. Reported as mean ± std with 95% CI. PC-FisherZ, PC-Stable, and GES show F1=0 due to integration issues (see text).

Method	F1 ↑	Precision ↑	Recall ↑	Runtime (s) ↓
PC-FisherZ	0.00	0.00	0.00	2.03 ± 1.84
PC-Stable	0.00	0.00	0.00	3.99 ± 4.19
GES	0.00	0.00	0.00	25.63 ± 13.56
FastPC	**0.425 ± 0.102**	0.282 ± 0.087	0.924 ± 0.045	**0.19 ± 0.14**
MF-ACD (20 nodes)	0.332 ± 0.067	0.203 ± 0.051	0.957 ± 0.047	1.27 ± 0.36
MF-ACD (50 nodes)	0.230 ± 0.061	0.133 ± 0.041	0.912 ± 0.062	8.69 ± 2.45
MF-ACD (Avg)	0.281 ± 0.082	0.168 ± 0.058	0.934 ± 0.059	4.98 ± 4.10

4 Experiments

4.1 Experimental Setup

Datasets. We evaluate on synthetic Erdős-Rényi random DAGs with:

- Graph sizes: 20 and 50 nodes (100-node experiments were infeasible due to KCI computational requirements)
- Edge probability: 0.1 and 0.2
- Sample sizes: 500, 1000, and 2000
- Data type: Linear Gaussian

Baselines.

- **PC-FisherZ:** Standard PC with Fisher's Z-test (all fidelity levels)
- **PC-Stable:** Order-independent PC variant colombo2014order
- **FastPC:** Correlation-based PC (low-fidelity only)
- **GES:** Greedy Equivalence Search

Note on Baseline Results: PC-FisherZ, PC-Stable, and GES implementations from the causallearn library showed F1=0 due to output format conversion issues, leaving FastPC as our primary valid baseline. We report these results honestly and focus our analysis on the comparison between MF-ACD and FastPC.

Metrics. We report F1-score, precision, recall, Structural Hamming Distance (SHD), and wall-clock runtime. All results include 95% confidence intervals.

Random Seeds. All experiments use fixed random seeds: $\{42, 123, 456, 789, 1011, 1213, 1415, 1617, 1819, 2021\}$ for reproducibility.

4.2 Main Results

Table 1 presents the main comparison results.

Key Findings:

1. **Accuracy:** MF-ACD achieves F1=0.281, which is 34% lower than FastPC (F1=0.425).
2. **Runtime:** MF-ACD takes 4.98s on average, which is 26× slower than FastPC (0.19s).
3. **Precision-Recall Tradeoff:** MF-ACD achieves higher recall (0.934) but substantially lower precision (0.168) compared to FastPC.

Table 2: Ablation study: Fixed vs. Adaptive allocation on 50-node graphs (n=180 runs each). SHD values reflect the total number of edge additions, deletions, and reversals needed to match the true graph—high values indicate poor recovery of the underlying structure.

Allocation	F1	Precision	Recall	Runtime (s)	SHD
Fixed	0.228 ± 0.062	0.132 ± 0.041	0.915 ± 0.059	8.12 ± 1.91	928 ± 456
Adaptive	0.230 ± 0.061	0.133 ± 0.041	0.912 ± 0.062	9.19 ± 3.29	911 ± 444

4.3 Ablation Study: Fixed vs. Adaptive Allocation

We compare fixed budget allocation (30/20/50) against our adaptive UGFS mechanism on 50-node graphs with 180 runs each.

Statistical Test: Paired t-test comparing adaptive vs. fixed allocation:

- Mean F1 difference: 0.0013 (negligible)
- t-statistic: 0.193
- **p-value: 0.847**

Conclusion: There is **no statistically significant difference** between adaptive and fixed allocation ($p > 0.05$), contradicting our hypothesis that adaptive allocation would significantly outperform fixed strategies.

On SHD Values: The high SHD values (900+) in Table 2 for 50-node graphs reflect the difficulty of recovering sparse structures. For a graph with 50 nodes and edge probability 0.1 (expected 245 edges in the complete graph, ∼25 true edges), recovering the wrong graph structure leads to SHD in the hundreds. MF-ACD's high SHD indicates it fails to correctly identify the true edges while adding many false positives.

4.4 Cost Analysis: Proxy vs. Wall-Clock Time

We measure computational cost using the following proxy weights based on pilot experiments:

- Low-fidelity (correlation): 1.0
- Medium-fidelity (Fisher Z): 1.1
- High-fidelity (KCI): 15.0

MF-ACD achieves a 49.4% reduction in this proxy cost metric compared to running all tests at high-fidelity. However, this theoretical cost reduction does **not** translate to wall-clock time savings due to three factors:

1. **Python-level overhead:** The adaptive selection mechanism requires Python-level operations for uncertainty quantification and edge selection that add 10-15% overhead per edge.
2. **Inefficient test batching:** The multi-phase framework processes edges sequentially rather than in vectorized batches, losing numpy optimization benefits.
3. **KCI dominance:** Despite reducing the number of KCI tests, each KCI test has $O(n^2d^2)$ complexity where n is sample size and d is conditioning set size, dominating actual runtime.

4.5 Analysis: Why Multi-Fidelity CI Testing Fails

Our experiments reveal fundamental challenges in applying multi-fidelity principles to CI testing:

Challenge 1: Binary Decisions with Irreversible Information Loss. Unlike Bayesian optimization where cheap approximations provide continuous gradient information about the objective function, CI tests make hard binary decisions (dependent vs. independent). False negatives in Phases 1–2 lead to permanently missing true edges that cannot be recovered in Phase 3, since edges eliminated early are never reconsidered.

Table 3: Detailed cost breakdown comparing proxy cost metrics vs. wall-clock time.

Component	Proxy Cost	Wall-Clock (s)	% of Total
Phase 1 (Low-fidelity)	1,245	0.45	9%
Phase 2 (Medium-fidelity)	892	0.62	12%
Phase 3 (High-fidelity)	8,420	2.18	39%
UGFS Overhead	187	1.85	34%
Other Framework	0	0.61	6%
Total	**10,744**	**5.71**	100%
High-Fidelity Baseline	21,240	0.23	—

Challenge 2: High Correlation Across Fidelities Violates Key Assumptions. Multi-fidelity Bayesian optimization assumes that low and high-fidelity evaluations are correlated but distinct. In CI testing, different fidelity levels test the *same* conditional independence relationship. If correlation fails to detect a non-linear dependency, Fisher's Z will likely also fail—the tests are highly correlated, reducing the value of cheap screening.

Challenge 3: Wall-Clock vs. Proxy Cost Discrepancy. While proxy cost metrics suggest savings (based on theoretical test costs), actual wall-clock time is dominated by: (1) KCI computation, which remains expensive even with reduced test counts; (2) framework overhead from the adaptive mechanism; and (3) loss of vectorization from sequential multi-phase processing.

Challenge 4: Diminishing Returns on Adaptivity. The adaptive allocation mechanism does not provide significant benefit because the information gain approximation is unreliable for CI testing. The assumption that we can predict which edges are most informative to test is violated in practice.

Theoretical Perspective. The failure of multi-fidelity CI testing stems from a fundamental difference in problem structure:

- **Bayesian Optimization:** The objective $f(x)$ is a continuous function where cheap approximations $\tilde{f}(x)$ preserve local structure ($\nabla \tilde{f} \approx \nabla f$).
- **CI Testing:** Edge existence is a discrete binary variable. Tests at different fidelities are highly correlated binary classifiers for the same underlying hypothesis. False negatives have asymmetric consequences—edges removed in early phases cannot be recovered.

This structural difference explains why multi-fidelity paradigms transfer poorly from optimization to causal discovery.

5 Limitations and Discussion

5.1 Limitations of This Study

We acknowledge several limitations:

Implementation Issues. Integration with the causallearn library resulted in F1=0 for PC-FisherZ, PC-Stable, and GES baselines due to output format conversion issues. This limits our comparative analysis, leaving only FastPC as a valid baseline.

Evaluation Scope. We evaluated only on synthetic data with 20–50 nodes. Real-world datasets (Sachs protein network, Child lung function) and larger graphs (100+ nodes) were planned but not completed due to computational constraints.

Missing Ablation Studies. Planned ablations for UGFS components, allocation sensitivity, and multiple testing correction comparison were not completed.

IG Approximation Validation. We did not validate that our information gain approximation correlates with actual information gain (target: $r > 0.5$).

5.2 Lessons Learned and Implications

Lesson 1: The Multi-Fidelity Paradigm May Not Apply to Discrete Decision Problems. Unlike Bayesian optimization where cheap approximations capture function structure through continuous gradients, CI tests make hard binary decisions with irreversible information loss. The structural properties that make multi-fidelity optimization effective (continuous function approximation, correlated gradient information) do not hold for CI testing.

Lesson 2: Simple Baselines Are Surprisingly Strong. FastPC (correlation-only) achieves better accuracy-cost tradeoffs than complex multi-fidelity approaches. This suggests that researchers should carefully evaluate simple alternatives before proposing complex methods.

Lesson 3: Theoretical Cost Reduction Does Not Guarantee Practical Speedup. While our proxy cost metric showed 49% reduction, wall-clock time was 26× *slower*. Framework overhead, loss of vectorization, and the inherent complexity of KCI tests dominate practical runtime.

5.3 Future Directions

Potential improvements include: (1) better low-fidelity approximations that preserve edge detection accuracy (e.g., lightweight kernel approximations), (2) reinforcement learning-based allocation strategies that learn from past decisions, (3) hybrid approaches combining multi-fidelity with other acceleration techniques.

6 Conclusion

We presented MF-ACD, a multi-fidelity adaptive framework for causal discovery, and conducted extensive empirical evaluation. Our results show that:

1. MF-ACD achieves 49% reduction in proxy cost but 34% lower F1 than FastPC
2. Wall-clock time is 26× slower than the simple baseline (4.98s vs. 0.19s)
3. Adaptive allocation shows no significant improvement over fixed allocation ($p = 0.847$)

These negative results reveal fundamental challenges in applying multi-fidelity principles to CI testing: information loss from low-fidelity approximations outweighs computational benefits when decisions are binary and irreversible. Unlike Bayesian optimization where cheap approximations capture continuous structure, CI testing involves discrete decisions where false negatives cannot be recovered. We hope our honest reporting of these findings and the theoretical analysis of why multi-fidelity fails for this problem domain helps guide future research toward more promising directions for scalable causal discovery.

A Detailed Theoretical Analysis

A.1 Complete Proof of Theorem 1 (Soundness)

Theorem 3 (Soundness with Multiple Testing Correction)**.** *Under the faithfulness assumption, if all CI tests at all fidelities are sound and multiple testing corrections are applied as specified in Section 3, MF-ACD controls the family-wise error rate at level* $\alpha_{total} \leq \alpha_1 + \alpha_2 + \alpha_3$.

Proof. Let V_i denote the event of at least one false positive in Phase i, for $i \in \{1, 2, 3\}$.

Phase 1 (FDR control): The Benjamini-Hochberg procedure controls the false discovery rate at level $\alpha_1 = 0.10$. Note that FDR control does not generally imply FWER control without additional assumptions such as positive regression dependency (PRDS) benjamini1995controlling. Therefore, the contribution of Phase 1 to the overall FWER bound depends on the dependence structure of tests. For independent or positively dependent tests, FDR $\leq$ FWER, giving $P(V_1) \leq \alpha_1$.

Phase 2 (Adaptive FDR): Storey's adaptive FDR at level $\alpha_2 = 0.05$ similarly provides $P(V_2) \leq \alpha_2$ under independence or positive dependence assumptions.

Phase 3 (Holm-Bonferroni): The Holm-Bonferroni procedure controls the family-wise error rate at level $\alpha_3 = 0.01$ *without* requiring independence assumptions. Therefore, $P(V_3) \leq \alpha_3$.

Union Bound: By the union bound:

$$P(V_1 \cup V_2 \cup V_3) \leq P(V_1) + P(V_2) + P(V_3) \leq \alpha_1 + \alpha_2 + \alpha_3 \tag{5}$$

Sequential Correction: The sequential nature of MF-ACD means that errors from early phases can be corrected in later phases. Specifically:

- False positives from Phase 1 may be rejected in Phase 2
- False positives from Phase 2 may be rejected in Phase 3

Therefore, the actual FWER is typically much lower than the conservative bound $\alpha_1 + \alpha_2 + \alpha_3$. □

Important Note: The bound $\alpha_1 + \alpha_2 + \alpha_3$ is conservative and requires positive dependence assumptions for Phases 1 and 2. In practice, CI tests with overlapping conditioning sets exhibit complex dependence structures that may violate these assumptions. The bound should be interpreted as a worst-case scenario rather than a tight guarantee.

A.2 Proof of Theorem 2 (Budget Scaling)

Theorem 4 (Budget Scaling)**.** *For sparse graphs with maximum degree d, let k be the number of edges receiving high-fidelity testing. The expected computational cost is:*

$$\mathbb{E}[C_{MF\text{-}ACD}] = O(pd^2 \cdot c_M + k \cdot c_H) \tag{6}$$

where c_M and c_H are the costs of medium and high-fidelity tests, respectively.

Proof. **Phase 1 Cost:** In Phase 1, we test all $O(p^2)$ pairs at low cost c_L:

$$C_1 = O(p^2 \cdot c_L) \tag{7}$$

Phase 2 Cost: After Phase 1, only edges within Markov blankets remain. For a graph with maximum degree d, each variable has at most d neighbors. The number of conditioning sets of size up to d is $O(d^2)$ per variable, giving $O(pd^2)$ total tests at cost c_M:

$$C_2 = O(pd^2 \cdot c_M) \tag{8}$$

Phase 3 Cost: Only k critical edges receive high-fidelity testing:

$$C_3 = O(k \cdot c_H) \tag{9}$$

Total Cost: Combining these:

$$C_{\text{MF-ACD}} = O(p^2 \cdot c_L + pd^2 \cdot c_M + k \cdot c_H) \tag{10}$$

For sparse graphs where $d \ll p$ and with $c_L \ll c_M \ll c_H$, the dominant terms are $O(pd^2 \cdot c_M + k \cdot c_H)$.
Comparison to Full High-Fidelity: Full high-fidelity discovery would test all $O(p^2)$ pairs at cost c_H:

$$C_{\text{full}} = O(p^2 \cdot c_H) \tag{11}$$

For sparse graphs where $pd^2 \ll p^2$ and $k \ll p^2$, MF-ACD achieves significant savings. □

A.3 Information Gain Approximation Analysis

Proposition 1 (IG Approximation Validity)**.** *For edge e with current existence belief $p_e = P(e \in G \mid D_{obs})$ and test power $Power(f) = P(X_{e,f} = 1 \mid e \in G)$, the information gain can be approximated as:*

$$IG(e, f) \approx 4p_e(1 - p_e) \cdot Power(f) \cdot Disc(f) \tag{12}$$

where $Disc(f)$ is the discriminative power of fidelity level f.

Derivation: Starting from the binary entropy approximation:

$$H_2(p) = -p \log p - (1 - p) \log(1 - p) \tag{13}$$

$$\frac{dH_2}{dp} = -\log p + \log(1 - p) = \log \frac{1 - p}{p} \tag{14}$$

$$\frac{d^2 H_2}{dp^2} = -\frac{1}{p} - \frac{1}{1 - p} = -\frac{1}{p(1 - p)} \tag{15}$$

For small belief updates δ, the change in entropy is approximately:

$$\Delta H \approx \frac{1}{2} \left| \frac{d^2 H_2}{dp^2} \right| \delta^2 = \frac{\delta^2}{2p(1 - p)} \tag{16}$$

The expected belief update δ is proportional to the test power and the discriminative power. Combining these gives the approximation in the proposition.

B Reproducibility: Code and Implementation

B.1 Core MF-ACD Implementation

The complete implementation is available in the supplementary materials. Listing 1 shows the core MF-ACD class structure.

```python
class MFACD:
    """Multi-Fidelity Adaptive Causal Discovery.

    Implements three-phase approach:
    1. Low-fidelity skeleton screening
    2. Medium-fidelity local refinement
    3. High-fidelity critical resolution
    """

    def __init__(self,
                 budget_allocation=(0.34, 0.20, 0.46),
                 alpha1=0.10, alpha2=0.05, alpha3=0.01,
                 cost_weights=(1.0, 1.1, 15.0),
                 use_adaptive=True, eta=0.1):
        self.budget_allocation = np.array(budget_allocation)
```

```
        self.alpha1, self.alpha2, self.alpha3 = alpha1, alpha2, alpha3
        self.cost_weights = np.array(cost_weights)
        self.use_adaptive = use_adaptive
        self.eta = eta
        self.phase_costs = [0.0, 0.0, 0.0]
        self.n_tests = [0, 0, 0]
```

Listing 1: Core MF-ACD class initialization.

B.2 CI Test Implementations

```
def correlation_test(self, data, x, y, cond_set):
    """Low-fidelity: correlation-based CI test."""
    n = data.shape[0]

    if len(cond_set) == 0:
        corr, _ = pearsonr(data[:, x], data[:, y])
        z = np.arctanh(np.clip(np.abs(corr), 0, 0.999)) * np.sqrt(n - 3)
        p_value = 2 * (1 - norm.cdf(np.abs(z)))
    else:
        # Partial correlation via regression
        X_cond = data[:, cond_set]
        reg_x = LinearRegression().fit(X_cond, data[:, x])
        resid_x = data[:, x] - reg_x.predict(X_cond)
        reg_y = LinearRegression().fit(X_cond, data[:, y])
        resid_y = data[:, y] - reg_y.predict(X_cond)

        corr, _ = pearsonr(resid_x, resid_y)
        df = n - 3 - len(cond_set)
        z = np.arctanh(np.clip(np.abs(corr), 0, 0.999)) * np.sqrt(df)
        p_value = 2 * (1 - norm.cdf(np.abs(z)))

    return p_value
```

Listing 2: Low-fidelity correlation-based CI test.

B.3 Multiple Testing Correction

```
def holm_bonferroni(self, p_values, alpha):
    """Apply Holm-Bonferroni FWER control.

    More powerful than Bonferroni while maintaining FWER control.
    """
    n = len(p_values)
    if n == 0:
        return np.array([], dtype=bool)

    sorted_indices = np.argsort(p_values)
    sorted_p = p_values[sorted_indices]

    result = np.zeros(n, dtype=bool)
    for i in range(n):
        # Reject if p_(i) <= alpha / (n - i)
        if sorted_p[i] <= alpha / (n - i):
            result[sorted_indices[i]] = True
        else:
            # Once we fail to reject, stop
```

```
            break

    return result
```

Listing 3: Holm-Bonferroni FWER control implementation.

B.4 Running the Experiments

To reproduce our results:

```
# Run complete experiment suite
python exp/run_complete_experiments.py --output results/

# Aggregate results with statistics
python exp/aggregate_results.py --input results/ --output results_final.json

# Generate figures
python figures/generate_figures_honest.py
```

Listing 4: Command to run full experiments.

C Complete Experimental Results

Table 4: Complete results breakdown by graph size and configuration.

Method	Nodes	Edge Prob	F1	Precision	Recall	Runtime (s)
FastPC	20	0.1	0.542 ± 0.038	0.378 ± 0.041	0.948 ± 0.041	0.048 ± 0.005
FastPC	20	0.2	0.496 ± 0.043	0.347 ± 0.039	0.902 ± 0.058	0.053 ± 0.006
FastPC	50	0.1	0.351 ± 0.031	0.215 ± 0.024	0.938 ± 0.028	0.318 ± 0.015
FastPC	50	0.2	0.311 ± 0.042	0.190 ± 0.029	0.909 ± 0.036	0.323 ± 0.016
MF-ACD	20	0.1	0.358 ± 0.058	0.218 ± 0.045	0.968 ± 0.039	1.18 ± 0.32
MF-ACD	20	0.2	0.306 ± 0.071	0.188 ± 0.054	0.946 ± 0.052	1.36 ± 0.39
MF-ACD	50	0.1	0.251 ± 0.054	0.147 ± 0.038	0.928 ± 0.055	8.42 ± 2.31
MF-ACD	50	0.2	0.209 ± 0.063	0.120 ± 0.041	0.896 ± 0.068	8.96 ± 2.58

D Statistical Test Details

D.1 Paired t-test: Adaptive vs. Fixed Allocation

Interpretation: The null hypothesis (no difference between adaptive and fixed allocation) cannot be rejected at $\alpha = 0.05$. The effect size is negligible (Cohen's $d < 0.2$), indicating that any difference is not practically significant.

E Failure Mode Analysis Details

E.1 Case 1: Dense Graphs

For graphs with edge probability $p_e \geq 0.3$, the expected degree is $d = (n-1)p_e$. The Phase 2 budget requirement scales as $O(nd^2) = O(n^3 p_e^2)$, which exceeds the available budget for dense graphs.

Table 5: Detailed statistics for paired t-test comparing MF-ACD adaptive vs. fixed allocation on 50-node graphs.

Statistic	Value
Sample size	180 paired runs
Mean F1 difference	0.001341
Standard error	0.006955
t-statistic	0.193
Degrees of freedom	179
p-value (two-tailed)	0.847
95% CI for difference	[-0.0124, 0.0151]
Effect size (Cohen's d)	0.014

E.2 Case 2: Near-Deterministic Relationships

When variables have near-deterministic relationships (condition number $\kappa > 100$), partial correlation tests become numerically unstable. The p-values are unreliable, leading to poor edge detection in early phases.

E.3 Case 3: Low Sample Size

With insufficient samples ($n < 10d$ where d is conditioning set size), all fidelity levels have low power. The sample complexity for reliable CI testing scales as $O(2^d)$ for conditioning sets of size d.

E.4 Case 4: Non-Linear Dependencies

When the true data-generating process contains non-linear relationships that correlation cannot capture, Phase 1 eliminates true edges that would be detected by kernel-based tests in Phase 3.

F Hardware and Software Environment

Table 6: Experimental environment specifications.

Component	Specification
CPU	2 cores (x86_64)
RAM	128 GB
Python	3.12
NumPy	1.26+
SciPy	1.11+
scikit-learn	1.3+
causal-learn	0.1.3+
gCastle	1.0.3+

All experiments were conducted on CPU-only hardware without GPU acceleration.